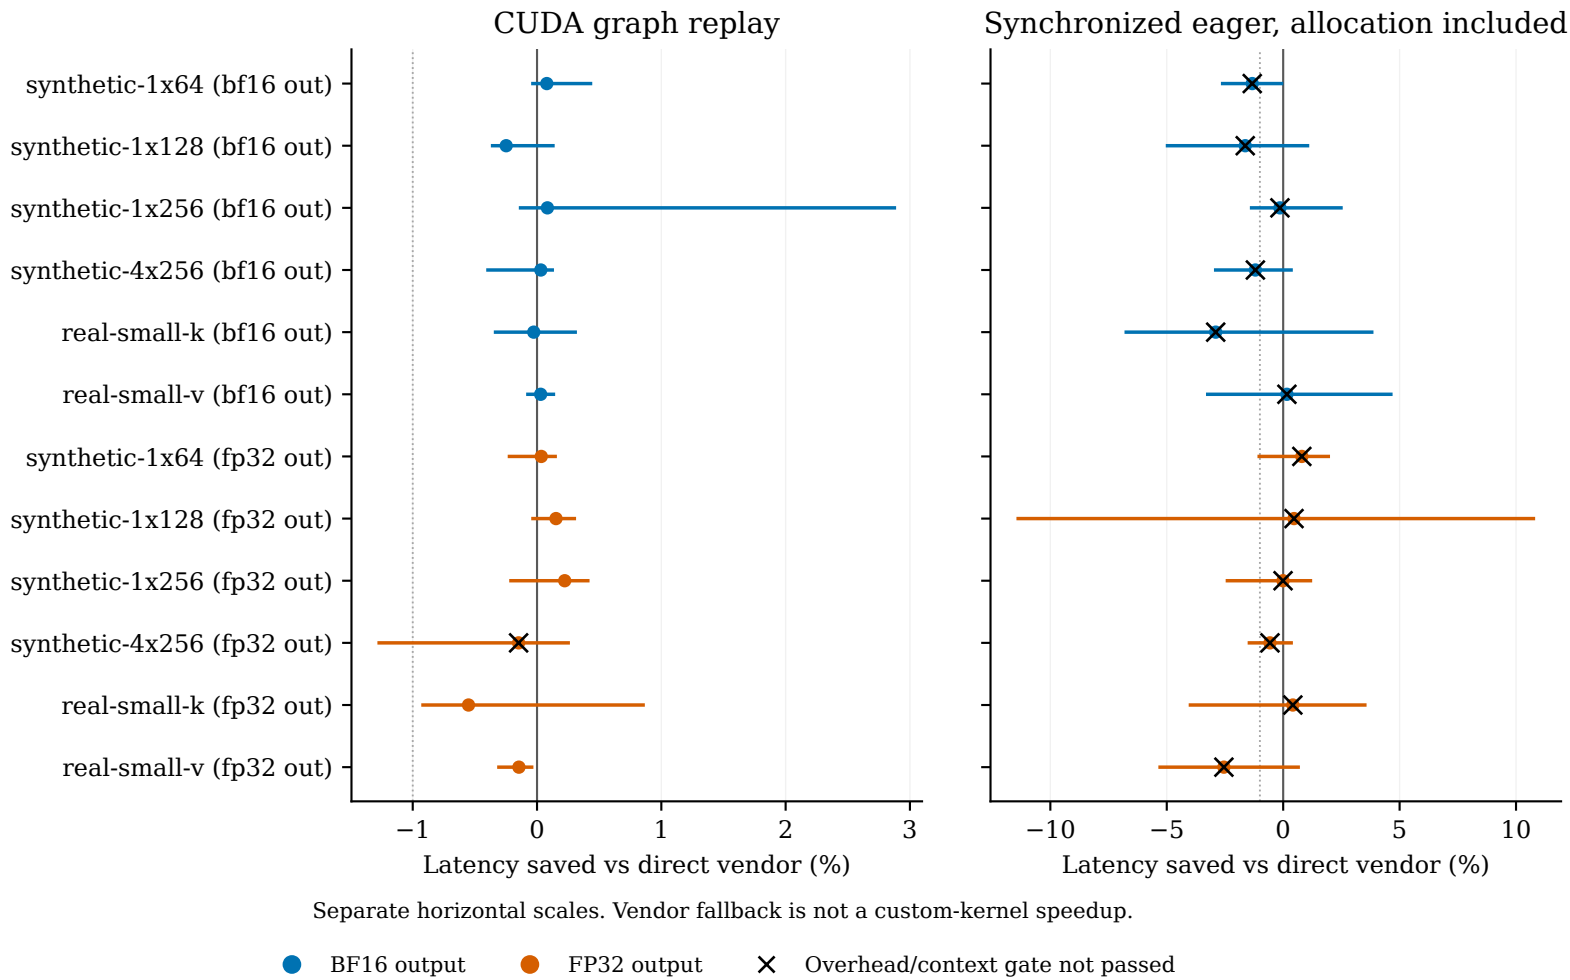
Prepared vendor fallback overhead
Native BF16 weights · FP32 accumulation · 31 paired rounds · stock GPU
CUDA graph replay
Synchronized eager, allocation included
synthetic-1x64 (bf16 out)
synthetic-1x128 (bf16 out)
synthetic-1x256 (bf16 out)
synthetic-4x256 (bf16 out)
real-small-k (bf16 out)
real-small-v (bf16 out)
synthetic-1x64 (fp32 out)
synthetic-1x128 (fp32 out)
synthetic-1x256 (fp32 out)
synthetic-4x256 (fp32 out)
real-small-k (fp32 out)
real-small-v (fp32 out)
−1
0
1
2
3
−10
−5
0
5
10
Latency saved vs direct vendor (%)
Latency saved vs direct vendor (%)
Separate horizontal scales. Vendor fallback is not a custom-kernel speedup.
BF16 output
FP32 output
Overhead/context gate not passed