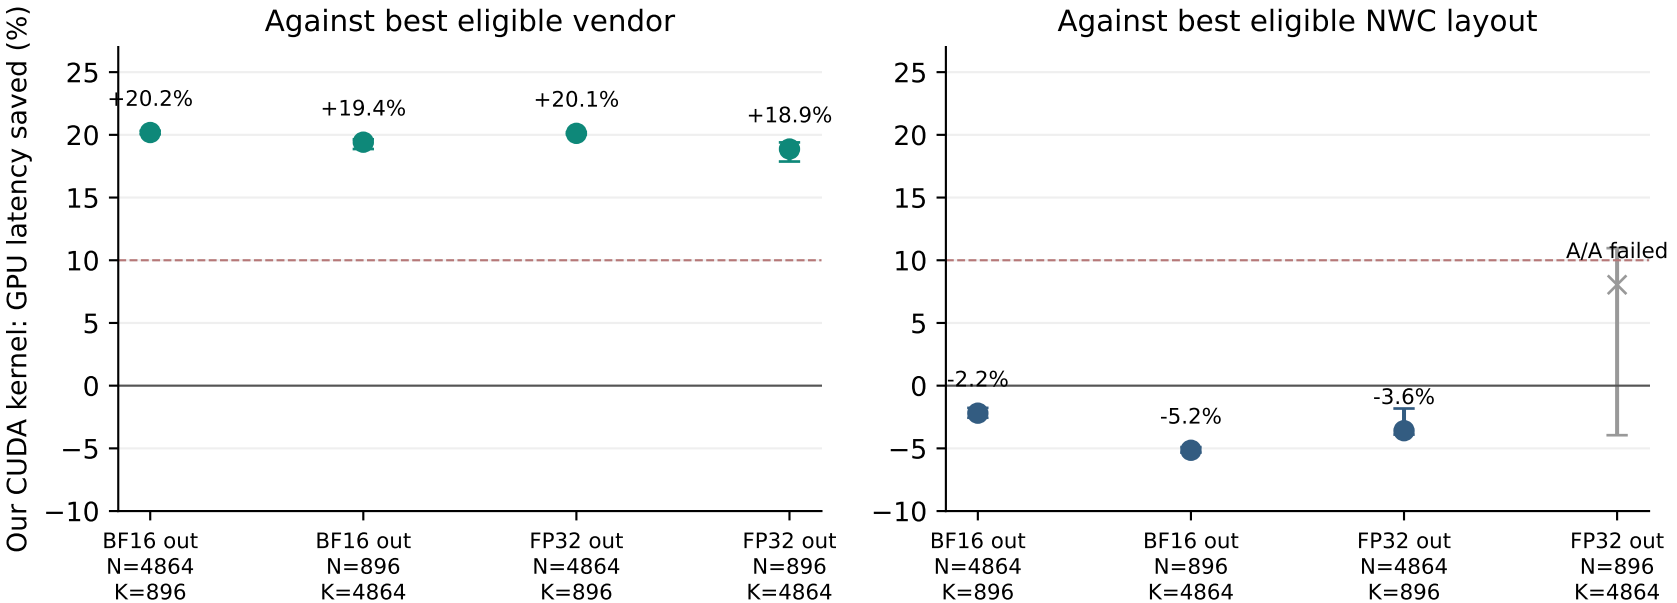
Real layer-0 weights, M=1, Performance profile
Against best eligible vendor
Against best eligible NWC layout
Our CUDA kernel: GPU latency saved (%)
25
20
15
10
5
0
−5
−10
+20.2%
+19.4%
+20.1%
+18.9%
-2.2%
-5.2%
-3.6%
A/A failed
BF16 out
N=4864
K=896
BF16 out
N=896
K=4864
FP32 out
N=4864
K=896
FP32 out
N=896
K=4864
Training only, seed 901. Exact BF16 weights and FP32 accumulation. Error bars: paired 95% intervals.